

Precision Sweep

Find quantisation-sensitive layers
and determine the per-layer precision



Baseline Error Profiling

Measure error under unified
scale/rounding/saturation



Per-layer Error Analysis

Locate where HW error grows and
adjust quantisation strategy



Per-layer Quantisation Strategy

Choose Q formats,
scales/rounding/saturation method



RTL Mapping

Implement quantisation module