

Precision Sweep

Find quantisation-sensitive layers



Baseline Error Profiling

Measure error under unified
scale/rounding



Per-layer Error Analysis

Locate where HW error grows and
adjust quantisation strategy



Per-layer Quantisation Strategy

Choose Q formats, scales/rounding
method



RTL Mapping

Implement quantisation module